

Tensorflow_Aquarium_Fish_Image_Classification

March 12, 2024

```
[2]: import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras import layers, models
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D,
    ↳AveragePooling2D, Concatenate, Flatten, Dense
from tensorflow.keras.layers import BatchNormalization, Activation, Add,
    ↳GlobalAveragePooling2D
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split, KFold
from sklearn.metrics import confusion_matrix
import math
import seaborn as sns
import numpy as np
import os
import cv2
import shutil
import pandas as pd
```

```
[3]: # silence warnings
import warnings
warnings.filterwarnings('ignore')
import absl.logging as absl_logging

# Set the logging level to suppress debug messages
tf.compat.v1.logging.set_verbosity(tf.compat.v1.logging.ERROR)
absl_logging.get_absl_logger().setLevel('ERROR')
```

```
[4]: # setup for multiple outputs from single cell
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = 'all'
```

The goal of this project is to construct an image classification system using a TensorFlow neural network to classify four common aquarium fish species. As a first step, 100 images were collected for each of the four species listed below. For each original image, 6 new transformed or processed versions were created:

- 1- Original image version (size: 300 x 300)
- 2- Original image version rotated 30 degrees counterclockwise (size: 300 x 300)
- 3- Original image version rotated 30 degrees clockwise (size: 300 x 300)
- 4- Horizontally flipped version (size: 300 x 300)
- 5- Horizontally flipped version rotated 30 degrees counterclockwise (size: 300 x 300)
- 6- Horizontally flipped version rotated 30 degrees clockwise (size: 300 x 300)

```
[142]: # Display the four fish species
# This script's ("Tensorflow_Aquarium_Fish_Image_Classification.ipynb")
# ↳ directory to provide
#   relative location of folder ('Pics') holding pictures
# Adjust this line to reflect any new location
location_of_this_ipynb_file = '/media/ijmg/SSD_FOUR_TB/ACADEMICS_101/
↳MY_PROJECTS/ADDED_PROJECTS/Fish/'

# Paths to images (in relative folder 'Pics') and their associated labels
path_to_barb_pic = os.path.join(location_of_this_ipynb_file, 'Pics/barb_001.
↳jpeg')
path_to_betta_pic = os.path.join(location_of_this_ipynb_file, 'Pics/betta_001.
↳jpeg')
path_to_gourami_pic = os.path.join(location_of_this_ipynb_file, 'Pics/
↳gourami_001.jpeg')
path_to_redtail_pic = os.path.join(location_of_this_ipynb_file, 'Pics/
↳redtail_001.jpeg')
image_paths = [path_to_barb_pic, path_to_betta_pic, path_to_gourami_pic,
↳path_to_redtail_pic]
labels = ['1-- Tiger Barb \n (Puntigrus tetrazona)',
          '2-- Siamese Fighting Fish - Crowntail \n (Betta splendens)',
          '3-- Three Spot Gourami \n (Trichopodus trichopterus)',
          '4-- Red-tailed Black Shark \n (Epalzeorhynchos bicolor)']

# Number of images to display in the grid
num_images = len(image_paths)

# Set up the subplot grid
num_rows = 2
num_cols = 2

# Create a figure and axes
fig, axes = plt.subplots(num_rows, num_cols, figsize=(5, 5))

# Iterate over the images and labels
for i, (image_path, label) in enumerate(zip(image_paths, labels)):
    # Load and display the image
    img = mpimg.imread(image_path)

    # Determine the subplot index
```

```

row_index = i // num_cols
col_index = i % num_cols

# Plot the image and set the title
axes[row_index, col_index].imshow(img)
axes[row_index, col_index].set_title(label)

# Remove axis ticks
axes[row_index, col_index].axis('off')

# Adjust layout for better visualization
plt.tight_layout()

# Show the plot
print('THE FOUR FISH SPECIES:\n')
plt.show();

```

THE FOUR FISH SPECIES:


1-- Tiger Barb
(*Puntigrus tetrazona*)



2-- Siamese Fighting Fish - Crowntail
(*Betta splendens*)



3-- Three Spot Gourami
(*Trichopodus trichopterus*)



4-- Red-tailed Black Shark
(*Epalzeorhynchos bicolor*)



```
[6]: # A function to create a variety of standardized, transformed training set
      ↪ images from a set of 100
      # original images
      # Function arguments: species name, list of image file names, input folder,
      ↪ output folder
      # Each original base input image will be replaced by 6 new transformed training
      ↪ images:
      # 1-- Original image version (size: 300 x 300)
      # 2-- Original image version rotated 30 degrees counterclockwise (size: 300 x
      ↪ 300)
      # 3-- Original image version rotated 30 degrees clockwise (size: 300 x 300)
      # 4-- Horizontally flipped version (size: 300 x 300)
      # 5-- Horizontally flipped version rotated 30 degrees counterclockwise (size:
      ↪ 300 x 300)
      # 6-- Horizontally flipped version rotated 30 degrees clockwise (size: 300 x
      ↪ 300)
      #
      # After utilizing the image_transformer function, each species will have
      ↪ approximately 600 training
      # set images
      def image_transformer(species_name, image_list, input_folder, output_folder):
          count = 0
          base_name = str(species_name)
          file_type = '.jpeg'
          for image in image_list:
              num_id = str(count).zfill(3)
              # extract image
              img = cv2.imread(input_folder + str(image))

              # resize image
              count = count + 1
              num_id = str(count).zfill(3)
              x_dim = 300
              y_dim = 300
              new_size = (x_dim, y_dim)
              img_resized = cv2.resize(img, new_size, interpolation= cv2.INTER_LINEAR)
              cv2.imwrite(os.path.join( output_folder + base_name + num_id +
              ↪ file_type ), img_resized)

              # flip and save resized image
              count = count + 1
              num_id = str(count).zfill(3)
              img_flip = cv2.flip(img_resized, 1)
              cv2.imwrite( os.path.join(output_folder + base_name + num_id +
              ↪ file_type), img_flip)
```

```

        # rotate both resized original and resized, flipped original +30 and
        ↪ -30 degrees then save images
        (h, w) = img_resized.shape[:2]
        center = (w / 2, h / 2)
        angle_1 = 30
        angle_2 = -30
        scale = 1
        M_1 = cv2.getRotationMatrix2D(center, angle_1, scale)
        img_rot_pos30 = cv2.warpAffine(img_resized, M_1, (w, h))
        count = count + 1
        num_id = str(count).zfill(3)
        cv2.imwrite( os.path.join(output_folder + base_name + num_id +
        ↪ file_type), img_rot_pos30)

        M_2 = cv2.getRotationMatrix2D(center, angle_2, scale)
        img_rot_neg30 = cv2.warpAffine(img_resized, M_2, (w, h))
        count = count + 1
        num_id = str(count).zfill(3)
        cv2.imwrite( os.path.join(output_folder + base_name + num_id +
        ↪ file_type), img_rot_neg30)

        M_1 = cv2.getRotationMatrix2D(center, angle_1, scale)
        img_rot_pos30 = cv2.warpAffine(img_flip, M_1, (w, h))
        count = count + 1
        num_id = str(count).zfill(3)
        cv2.imwrite( os.path.join(output_folder + base_name + num_id +
        ↪ file_type), img_rot_pos30)

        M_2 = cv2.getRotationMatrix2D(center, angle_2, scale)
        img_rot_neg30 = cv2.warpAffine(img_flip, M_2, (w, h))
        count = count + 1
        num_id = str(count).zfill(3)
        cv2.imwrite( os.path.join(output_folder + base_name + num_id +
        ↪ file_type), img_rot_neg30)

```

```

[7]: # A function to create directories if they do not exist
      # This function will create directories holding the transformed images
      def create_directory(directory_path):
          if not os.path.exists(directory_path):
              os.makedirs(directory_path)
          else:
              pass

```

```

[8]: # OBTAIN LIST OF ORIGINAL IMAGES
      # Each species of fish has its own folder holding original images

```

```

# Folders are located relative to this file
↳("Tensorflow_Aquarium_Fish_Image_Classification.ipynb")
barb_original_folder = os.path.join(location_of_this_ipynb_file,
↳'original_images/original_barbs/')
betta_original_folder = os.path.join(location_of_this_ipynb_file,
↳'original_images/original_bettas/')
gourami_original_folder = os.path.join(location_of_this_ipynb_file,
↳'original_images/original_gouramis/')
redtail_original_folder = os.path.join(location_of_this_ipynb_file,
↳'original_images/original_redtails/')
# Read original image names for each species into a unique image list
barb_original_image_list = os.listdir(barb_original_folder)
betta_original_image_list = os.listdir(betta_original_folder)
gourami_original_image_list = os.listdir(gourami_original_folder)
redtail_original_image_list = os.listdir(redtail_original_folder)

```

```

[9]: # PREPARE TRANSFORMED IMAGES
# Each species of fish has its own folder holding transformed images
# Folders are located relative to this file
↳("Tensorflow_Aquarium_Fish_Image_Classification.ipynb")
barb_transformed_folder = os.path.join(location_of_this_ipynb_file,
↳'transformed_images/barbs/')
betta_transformed_folder = os.path.join(location_of_this_ipynb_file,
↳'transformed_images/bettas/')
gourami_transformed_folder = os.path.join(location_of_this_ipynb_file,
↳'transformed_images/gouramis/')
redtail_transformed_folder = os.path.join(location_of_this_ipynb_file,
↳'transformed_images/redtails/')
create_directory(barb_transformed_folder)
create_directory(betta_transformed_folder)
create_directory(gourami_transformed_folder)
create_directory(redtail_transformed_folder)
# Send each species original image list to the image_transformer function to
↳make new image training versions
# For each function call, four parameters are passed:
# 1.) species base name for naming sequential images (i.e. "barb001.jpeg",
↳"barb002.jpeg" ... etc.),
# 2.) base input image list of original 100 base input images,
# 3.) input directory holding original 100 base input images for each species,
↳and
# 4.) output directory where expanded output training images will be written
↳for each species
image_transformer('barbs_', barb_original_image_list, barb_original_folder,
↳barb_transformed_folder)
image_transformer('bettas_', betta_original_image_list, betta_original_folder,
↳betta_transformed_folder)

```

```

image_transformer('gouramis_', gourami_original_image_list,
    ↳gourami_original_folder, gourami_transformed_folder)
image_transformer('redtails_', redtail_original_image_list,
    ↳redtail_original_folder, redtail_transformed_folder)
# This resulted in roughly 600 images for each of the fish species.
# These images are in the transformed_images folder

```

```

[104]: # The 600 images for each fish species was then divided between a training set
    ↳(roughly 80%)
    # and a test set (roughly 20%).
    # Specify path to dataset
dataset_path = location_of_this_ipynb_file + 'transformed_images/'

# Create ImageDataGenerator for splitting dataset
# ImageDataGenerator was used to add even more variety to the training set
    ↳images
data_generator = ImageDataGenerator(
rescale=1./255,          # Normalize pixel values to [0, 1]
validation_split=0.20) # 20% for validation (test) set

train_generator = data_generator.flow_from_directory(
    dataset_path,
    shuffle=True,
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical',
    subset='training'
)

test_generator = data_generator.flow_from_directory(
    dataset_path,
    shuffle=False,
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical',
    subset='validation'
)

```

Found 1920 images belonging to 4 classes.

Found 480 images belonging to 4 classes.

```

[11]: # Show the four class labels
print("Train set class labels:", train_generator.class_indices)
print("Test set class labels:", test_generator.class_indices)

# Show the class indices, class labels and class values

```

```

class_indices = train_generator.class_indices
class_labels = list(class_indices.keys())
class_values = list(class_indices.values())

print(class_indices)
print(class_labels)
print(class_values)

```

Train set class labels: {'barbs': 0, 'bettas': 1, 'gouramis': 2, 'redtails': 3}
 Test set class labels: {'barbs': 0, 'bettas': 1, 'gouramis': 2, 'redtails': 3}
 {'barbs': 0, 'bettas': 1, 'gouramis': 2, 'redtails': 3}
 ['barbs', 'bettas', 'gouramis', 'redtails']
 [0, 1, 2, 3]

```

[12]: # Count the number of images for each class
      # Dictionaries to hold class counts
      image_count_per_train_class = {label: 0 for label in class_labels}
      image_count_per_test_class = {label: 0 for label in class_labels}

      # Load each dictionary with class counts
      for class_label in train_generator.classes:
          class_name = class_labels[class_label]
          image_count_per_train_class[class_name] += 1

      for class_label in test_generator.classes:
          class_name = class_labels[class_label]
          image_count_per_test_class[class_name] += 1

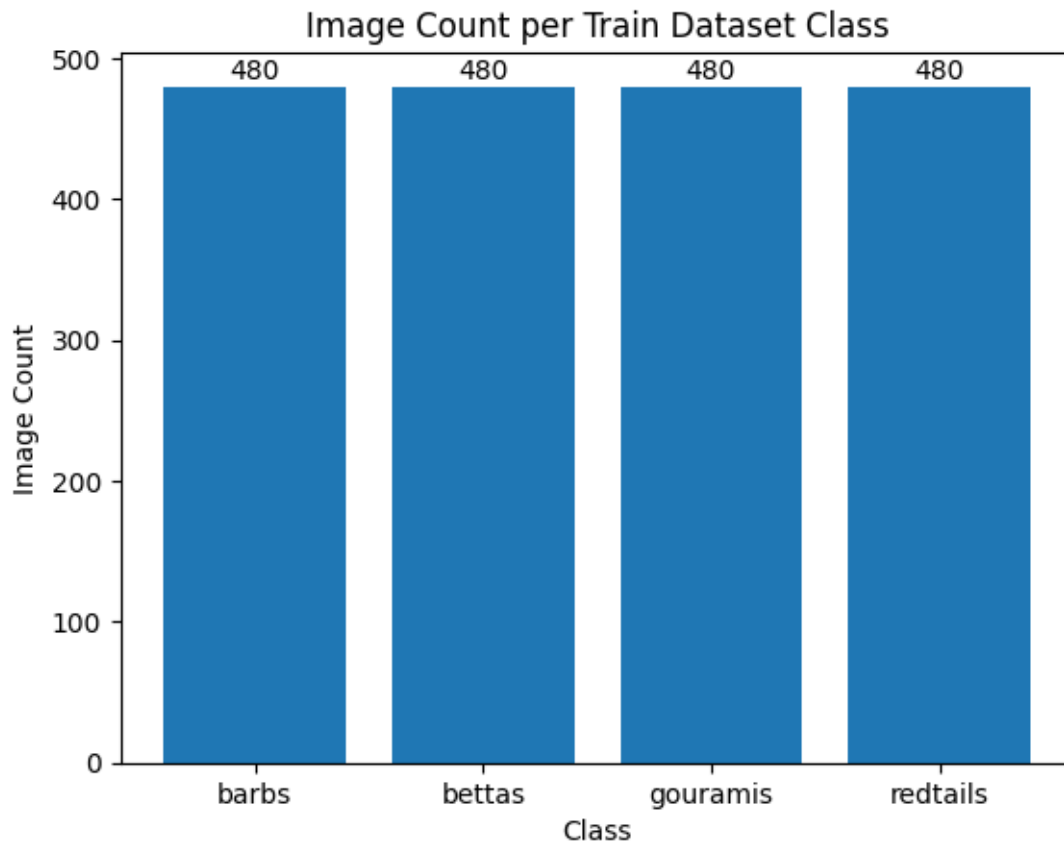
      # Print the image count for each class
      print('\nTraining Image Class Counts:')
      for class_name, count in image_count_per_train_class.items():
          print(f'{class_name}: {count} images')
      print('\nTest Image Class Counts:')
      for class_name, count in image_count_per_test_class.items():
          print(f'{class_name}: {count} images')

```

Training Image Class Counts:
 barbs: 480 images
 bettas: 480 images
 gouramis: 480 images
 redtails: 480 images

Test Image Class Counts:
 barbs: 120 images
 bettas: 120 images
 gouramis: 120 images
 redtails: 120 images


```
[13]: # Image count in each train dataset class
# Plot the bar chart of image count in each train class
fig, ax = plt.subplots()
ax.bar(image_count_per_train_class.keys(), image_count_per_train_class.values())
# Add numeric values at the top of each bar
for i, value in enumerate(image_count_per_train_class.values()):
    plt.text(i, value + 1, str(value), ha='center', va='bottom')
ax.set_xlabel('Class')
ax.set_ylabel('Image Count')
ax.set_title('Image Count per Train Dataset Class')
plt.show();
```

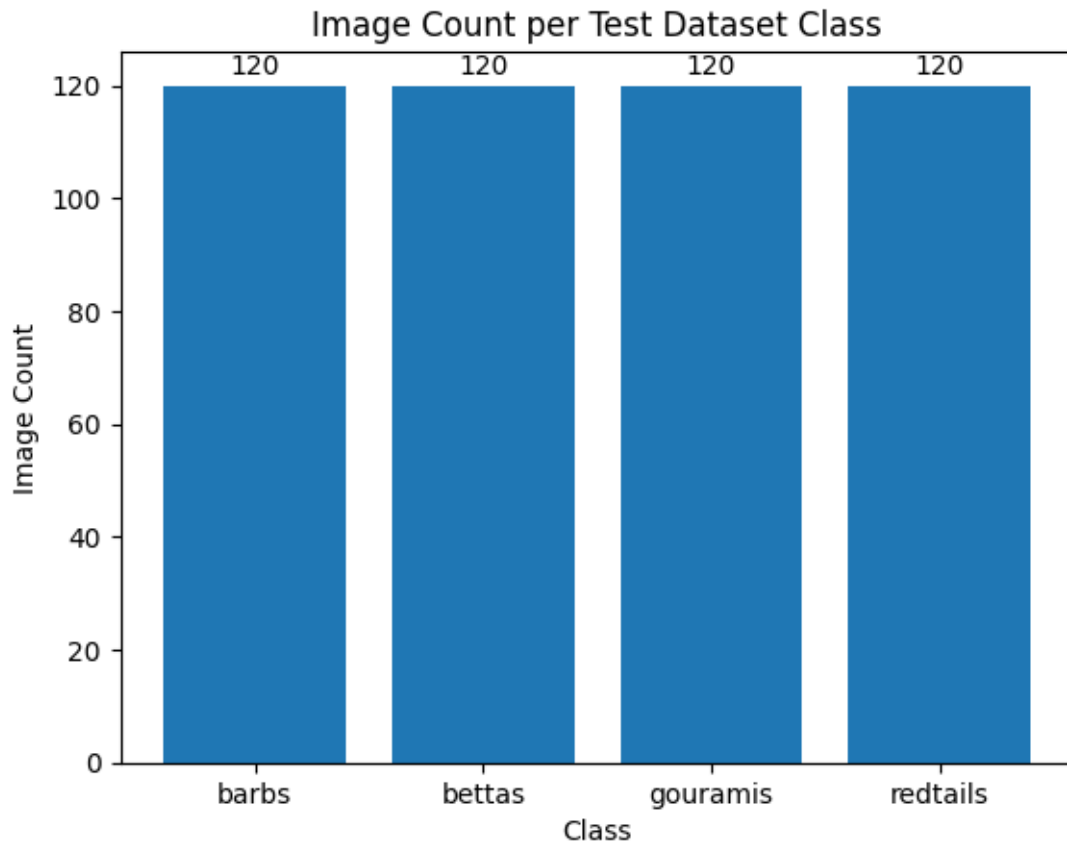


```
[14]: # Image count in each test dataset class
# Plot the bar chart of image count in each test class (slightly different
    ↪ method)
categories = image_count_per_test_class.keys()
values = image_count_per_test_class.values()
plt.bar(categories, values)
# Add numeric values at the top of each bar
for i, value in enumerate(image_count_per_test_class.values()):
```

```

plt.text(i, value + 1, str(value), ha='center', va='bottom')
# Add labels and title
plt.xlabel('Class')
plt.ylabel('Image Count')
plt.title('Image Count per Test Dataset Class')
# Display the plot
plt.show();

```



```

[15]: # Display sample of 15 random images from the training dataset
fig, axes = plt.subplots(3, 5, figsize=(15, 6))

for i in range(15):
    random_image, label = train_generator.next()
    class_name = list(train_generator.class_indices.keys())[np.argmax(label)]
    axes[i//5, i%5].imshow(random_image[0])
    axes[i//5, i%5].set_title(f'Training Class: {class_name}')
    axes[i//5, i%5].axis('off')

plt.show();

```



```
[16]: # Display sample of 15 random images from the test dataset
fig, axes = plt.subplots(3, 5, figsize=(15, 6))

for i in range(15):
    random_image, label = test_generator.next()
    class_name = list(test_generator.class_indices.keys())[np.argmax(label)]
    axes[i//5, i%5].imshow(random_image[0])
    axes[i//5, i%5].set_title(f'Test Class: {class_name}')
    axes[i//5, i%5].axis('off')

plt.show();
```



```
[17]: # Two different models will be compared:
# -- a general tensorflow model
```

```
# -- a Googlenet tensorflow model
# Each model will be tested with and without crossfold validation
# The model with the best performance (accuracy) will be reported more fully in
↳ terms of
#     number of images mislabeled and display of mislabeled images
```

```
[18]: # -----
# -----
# IMAGE CLASSIFICATION WITH GENERAL TENSORFLOW MODEL
# -----
# -----
```

```
[141]: # Function to define general model
def build_model(num_classes):
    model = tf.keras.Sequential([
        tf.keras.layers.Conv2D(filters=8, kernel_size=(3, 3), activation='relu',
                                input_shape=(224, 224, 3)),
        tf.keras.layers.MaxPooling2D(pool_size=(2, 2), strides=2),
        tf.keras.layers.Conv2D(filters=16, kernel_size=(3, 3), activation='relu'),
        tf.keras.layers.MaxPooling2D(pool_size=(2, 2), strides=2),
        tf.keras.layers.Conv2D(filters=32, kernel_size=(3, 3), activation='relu'),
        tf.keras.layers.MaxPooling2D(pool_size=(2, 2), strides=2),
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dense(64, activation='relu'),
        tf.keras.layers.Dense(num_classes, activation='softmax')
    ])
    return model
```

```
[21]: # Define number of classes
Num_classes = 4
# Create the general model
general_model = build_model(Num_classes)
```

```
[22]: # Compile Model
general_model.compile(
    optimizer=tf.keras.optimizers.SGD(learning_rate = 1e-3, momentum=0.9),
    loss='categorical_crossentropy',
    metrics=['accuracy'])
# Train the model
general_model_history = general_model.fit(train_generator, epochs=20 )
```

Epoch 1/20

```
2024-03-06 11:25:39.516062: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
[[{{node Placeholder/_0}}]]
```

60/60 [=====] - 38s 619ms/step - loss: 1.2631 -
accuracy: 0.4120
Epoch 2/20
60/60 [=====] - 38s 633ms/step - loss: 0.9854 -
accuracy: 0.5932
Epoch 3/20
60/60 [=====] - 41s 676ms/step - loss: 0.7839 -
accuracy: 0.6740
Epoch 4/20
60/60 [=====] - 40s 662ms/step - loss: 0.6456 -
accuracy: 0.7521
Epoch 5/20
60/60 [=====] - 40s 673ms/step - loss: 0.5627 -
accuracy: 0.7797
Epoch 6/20
60/60 [=====] - 39s 650ms/step - loss: 0.4561 -
accuracy: 0.8323
Epoch 7/20
60/60 [=====] - 39s 643ms/step - loss: 0.3831 -
accuracy: 0.8604
Epoch 8/20
60/60 [=====] - 40s 667ms/step - loss: 0.3240 -
accuracy: 0.8833
Epoch 9/20
60/60 [=====] - 39s 643ms/step - loss: 0.2820 -
accuracy: 0.9104
Epoch 10/20
60/60 [=====] - 42s 696ms/step - loss: 0.2434 -
accuracy: 0.9172
Epoch 11/20
60/60 [=====] - 41s 685ms/step - loss: 0.2207 -
accuracy: 0.9286
Epoch 12/20
60/60 [=====] - 42s 700ms/step - loss: 0.1389 -
accuracy: 0.9547
Epoch 13/20
60/60 [=====] - 42s 696ms/step - loss: 0.1236 -
accuracy: 0.9672
Epoch 14/20
60/60 [=====] - 45s 739ms/step - loss: 0.0925 -
accuracy: 0.9766
Epoch 15/20
60/60 [=====] - 44s 718ms/step - loss: 0.0778 -
accuracy: 0.9823
Epoch 16/20
60/60 [=====] - 43s 712ms/step - loss: 0.0648 -
accuracy: 0.9828
Epoch 17/20

```

60/60 [=====] - 44s 732ms/step - loss: 0.0515 -
accuracy: 0.9891
Epoch 18/20
60/60 [=====] - 46s 764ms/step - loss: 0.0314 -
accuracy: 0.9943
Epoch 19/20
60/60 [=====] - 47s 782ms/step - loss: 0.0210 -
accuracy: 0.9984
Epoch 20/20
60/60 [=====] - 37s 620ms/step - loss: 0.0151 -
accuracy: 0.9995

```


```

[72]: # Plot the training set loss
plt.figure(figsize=(10, 4))
plt.subplot(1, 2, 1)
plt.plot(general_model_history.history['loss'], label='Training Loss -- General Model')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training Loss -- General Model')
plt.xticks(np.linspace(0, 20, 5))
plt.legend()

# Plot the training set accuracy
plt.subplot(1, 2, 2)
plt.plot(general_model_history.history['accuracy'], label='Training Accuracy -- General Model')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Training Accuracy -- General Model')
plt.xticks(np.linspace(0, 20, 5))
plt.legend()

plt.tight_layout();
plt.show();

```



```
[24]: # Evaluate the model on the test set
general_model_predictions = general_model.predict(test_generator)
```


```
2024-03-06 11:43:07.848742: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
```

```
[[{{node Placeholder/_0}}]]
```

```
15/15 [=====] - 2s 131ms/step
```

```
[105]: # Create confusion matrix
general_model_predicted_labels = np.argmax(general_model_predictions, axis=1)
true_labels = test_generator.classes
conf_matrix = confusion_matrix(true_labels, general_model_predicted_labels)

# Plot the confusion matrix
plt.figure(figsize=(5, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=test_generator.class_indices, yticklabels=test_generator.
            class_indices)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix -- General Model')
plt.show();
```



```
[26]: # Count mislabeled images
# Get class names from folder names
class_names = sorted(os.listdir(dataset_path))

# Get true labels
true_labels = test_generator.classes

# Get predicted labels
general_model_predicted_labels = np.argmax(general_model_predictions, axis=1)

# Find mislabeled indices
general_mislabeled_indices = np.where(true_labels !=
    ↪ general_model_predicted_labels)[0]

print("Number of mislabeled images:", len(general_mislabeled_indices))
print("Number of predictions:", len(general_model_predictions))
print("Overall accuracy:", round((1- (len(general_mislabeled_indices)/
    ↪ len(general_model_predictions)))*100, 2), "percent" )
```

```
Number of mislabeled images: 76
Number of predictions: 480
Overall accuracy: 84.17 percent
```



```
[27]: # -----
# -----
# IMAGE CLASSIFICATION WITH GOOGLNET TENSORFLOW MODEL
# -----
# -----
```

```
[140]: # Functions to define Googlenet model
def inception_module(x, filters):
    # 1x1 convolution
    conv1x1_1 = Conv2D(filters[0], (1, 1), padding='same', activation='relu')(x)

    # 1x1 convolution followed by 3x3 convolution
    conv1x1_2 = Conv2D(filters[1], (1, 1), padding='same', activation='relu')(x)
    conv3x3 = Conv2D(filters[2], (3, 3), padding='same',
    ↪activation='relu')(conv1x1_2)

    # 1x1 convolution followed by 5x5 convolution
    conv1x1_3 = Conv2D(filters[3], (1, 1), padding='same', activation='relu')(x)
    conv5x5 = Conv2D(filters[4], (5, 5), padding='same',
    ↪activation='relu')(conv1x1_3)

    # 3x3 max pooling followed by 1x1 convolution
    maxpool = MaxPooling2D((3, 3), strides=(1, 1), padding='same')(x)
    conv1x1_4 = Conv2D(filters[5], (1, 1), padding='same',
    ↪activation='relu')(maxpool)

    # Concatenate all branches
    inception = Concatenate(axis=-1)([conv1x1_1, conv3x3, conv5x5, conv1x1_4])
    return inception

def googlenet():
    input_layer = Input(shape=(224, 224, 3))

    # Initial convolution layer
    x = Conv2D(64, (7, 7), strides=(2, 2), padding='same',
    ↪activation='relu')(input_layer)
    x = MaxPooling2D((3, 3), strides=(2, 2), padding='same')(x)

    # Inception modules
    x = inception_module(x, [64, 128, 128, 32, 32, 32])
    x = inception_module(x, [128, 192, 96, 64, 64, 64])

    # Additional inception modules...

    # Fully connected layers
    x = AveragePooling2D((7, 7))(x)
    x = Flatten()(x)
```

```

x = Dense(4, activation='softmax')(x)

model = tf.keras.Model(inputs=input_layer, outputs=x, name='googlenet')
return model

```

```

[29]: # Create Googlenet model
google_model = googlenet()

```

```

[30]: # Compile model
google_model.compile(
    optimizer=tf.keras.optimizers.SGD(learning_rate = 1e-3, momentum=0.9),
    loss='categorical_crossentropy',
    metrics=['accuracy'])
# Train the model
google_history = google_model.fit(train_generator, epochs=20 )

```

Epoch 1/20

2024-03-06 11:45:26.970374: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID_ARGUMENT: You must feed a value for placeholder tensor 'Placeholder/_0' with dtype int32

[[{{node Placeholder/_0}}]]

60/60 [=====] - 475s 8s/step - loss: 1.3274 - accuracy: 0.3620

Epoch 2/20

60/60 [=====] - 459s 8s/step - loss: 1.0719 - accuracy: 0.5703

Epoch 3/20

60/60 [=====] - 470s 8s/step - loss: 0.9076 - accuracy: 0.6427

Epoch 4/20

60/60 [=====] - 474s 8s/step - loss: 0.7956 - accuracy: 0.6932

Epoch 5/20

60/60 [=====] - 480s 8s/step - loss: 0.6950 - accuracy: 0.7349

Epoch 6/20

60/60 [=====] - 480s 8s/step - loss: 0.5862 - accuracy: 0.7839

Epoch 7/20

60/60 [=====] - 486s 8s/step - loss: 0.5238 - accuracy: 0.8151

Epoch 8/20

60/60 [=====] - 486s 8s/step - loss: 0.4792 - accuracy: 0.8349

Epoch 9/20

```

60/60 [=====] - 472s 8s/step - loss: 0.4353 - accuracy:
0.8573
Epoch 10/20
60/60 [=====] - 483s 8s/step - loss: 0.3764 - accuracy:
0.8682
Epoch 11/20
60/60 [=====] - 481s 8s/step - loss: 0.3519 - accuracy:
0.8750
Epoch 12/20
60/60 [=====] - 488s 8s/step - loss: 0.3422 - accuracy:
0.8776
Epoch 13/20
60/60 [=====] - 491s 8s/step - loss: 0.2926 - accuracy:
0.8953
Epoch 14/20
60/60 [=====] - 503s 8s/step - loss: 0.2711 - accuracy:
0.9094
Epoch 15/20
60/60 [=====] - 534s 9s/step - loss: 0.2605 - accuracy:
0.9104
Epoch 16/20
60/60 [=====] - 620s 10s/step - loss: 0.2730 -
accuracy: 0.9016
Epoch 17/20
60/60 [=====] - 549s 9s/step - loss: 0.2356 - accuracy:
0.9229
Epoch 18/20
60/60 [=====] - 503s 8s/step - loss: 0.2286 - accuracy:
0.9208
Epoch 19/20
60/60 [=====] - 485s 8s/step - loss: 0.2056 - accuracy:
0.9297
Epoch 20/20
60/60 [=====] - 481s 8s/step - loss: 0.1961 - accuracy:
0.9344

```

```

[73]: # Plot the training set loss
plt.figure(figsize=(10, 4))
plt.subplot(1, 2, 1)
plt.plot(google_history.history['loss'], label='Training Loss -- GoogleNet_
↳Model')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training Loss -- GoogleNet Model')
plt.xticks(np.linspace(0, 20, 5))
plt.legend()

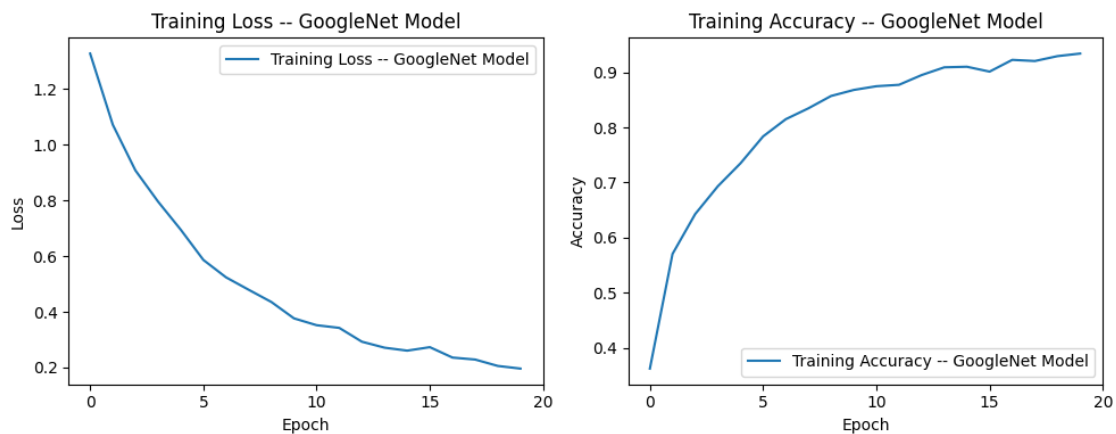
```

```

# Plot the training set accuracy
plt.subplot(1, 2, 2)
plt.plot(google_history.history['accuracy'], label='Training Accuracy -- GoogleNet Model')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Training Accuracy -- GoogleNet Model')
plt.xticks(np.linspace(0, 20, 5))
plt.legend()

plt.tight_layout();
plt.show();

```



```

[32]: # Evaluate the model on the test set
google_model_predictions = google_model.predict(test_generator)

```

```

2024-03-06 14:47:01.027631: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
[[{{node Placeholder/_0}}]]

```

```

15/15 [=====] - 36s 2s/step

```

```

[106]: # Create confusion matrix
google_model_predicted_labels = np.argmax(google_model_predictions, axis=1)
true_labels = test_generator.classes
conf_matrix = confusion_matrix(true_labels, google_model_predicted_labels)


# Plot the confusion matrix
plt.figure(figsize=(5, 4))

```

```

sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Reds',
            xticklabels=test_generator.class_indices, yticklabels=test_generator.
            class_indices)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix -- GoogleNet Model')
plt.show();

```



```

[34]: # Count mislabeled images
# Get class names from folder names
class_names = sorted(os.listdir(dataset_path))

# Get true labels
true_labels = test_generator.classes

# Get predicted labels
google_model_predicted_labels = np.argmax(google_model_predictions, axis=1)

# Find mislabeled indices
google_model_mislabeled_indices = np.where(true_labels !=
    google_model_predicted_labels)[0]

print("Number of mislabeled images:", len(google_model_mislabeled_indices))

```

```

print("Number of predictions:", len(google_model_predictions))
print("Overall accuracy:", round((1- (len(google_model_mislabeled_indices)/
↳len(google_model_predictions)))*100, 2), "percent" )

```

Number of mislabeled images: 79

Number of predictions: 480

Overall accuracy: 83.54 percent

```

[107]: #
↳
#
↳
# SUMMARY OF PERFORMANCES FOR NON CROSSFOLD VALIDATION (CFV) GENERAL AND
↳GOOGLNET MODEL VARIATIONS
#
↳
#
↳
↳

# General Model Confusion Matrix
plt.figure(figsize=(10, 4))
plt.subplot(1, 2, 1)
general_model_predicted_labels = np.argmax(general_model_predictions, axis=1)
true_labels = test_generator.classes
conf_matrix = confusion_matrix(true_labels, general_model_predicted_labels)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
↳xticklabels=test_generator.class_indices, yticklabels=test_generator.
↳class_indices)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix -- General Model')


# GoogleNet Model Confusion Matrix
plt.subplot(1, 2, 2)
google_model_predicted_labels = np.argmax(google_model_predictions, axis=1)
true_labels = test_generator.classes
conf_matrix = confusion_matrix(true_labels, google_model_predicted_labels)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Reds',
↳xticklabels=test_generator.class_indices, yticklabels=test_generator.
↳class_indices)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix -- GoogleNet Model')
plt.tight_layout();
plt.show();

```

```

general_accuracy = str(round((1- (len(general_mislabeled_indices)/
↳len(general_model_predictions)))*100, 2)) + ' %'
google_accuracy = str(round((1- (len(google_model_mislabeled_indices)/
↳len(google_model_predictions)))*100, 2)) + ' %'
# Organize Data
data = {'Model': ['General', 'GoogleNet'],
        'Predictions': [len(general_model_predictions),
↳len(google_model_predictions)],
        'Errors': [len(general_mislabeled_indices),
↳len(google_model_mislabeled_indices)],
        'Accuracy': [general_accuracy, google_accuracy]}
# Construct a DataFrame
df = pd.DataFrame.from_dict(data).set_index('Model')
# Display the DataFrame
print(df);

```



	Predictions	Errors	Accuracy
Model			
General	480	76	84.17 %
GoogleNet	480	79	83.54 %

OBSERVATIONS FOR NON CROSSFOLD VALIDATION (CFV) MODEL VARIATIONS:

- 1- Each model may be suffering from overfitting
- 2- Both models had difficulty separating bettas from gouramis
- 3- The general model also had trouble differentiating barbs from bettas (12 errors)
- 4- The googlenet model also had trouble differentiating redtails from barbs (13 errors)
- 5- For both models, each row contains 120 image predictions as expected

```
[127]: # -----
# -----
# TRY CROSSFOLD VALIDATION (CFV) MODEL VARIATIONS
# -----
# -----

# List all classes (subdirectories) in the root directory
classes = [class_name for class_name in os.listdir(dataset_path) if os.path.
↳ isdir(os.path.join(dataset_path, class_name))]

# Define lists to store file paths and associated labels to set up cross_
↳ validation
file_paths = []
labels = []

# Iterate over classes
for class_name in classes:
    class_path = os.path.join(dataset_path, class_name)

    # List all files in the class directory
    class_files = [os.path.join(class_path, file) for file in os.
↳ listdir(class_path) if file.endswith('.jpeg')]

    # Append file paths and labels to the lists
    file_paths.extend(class_files)
    labels.extend([class_name] * len(class_files))

# Create a DataFrame to store file paths and labels
df = pd.DataFrame({'Filepath': file_paths, 'Label': labels})

# Use train_test_split to split the DataFrame into training, validation, and_
↳ test sets
# First, split the dataset: 80% into training and 20% into testing sets
train_df, test_df = train_test_split(df, test_size=0.2, random_state=42)
# Then, further split the training set: 80% into training set and 20% into_
↳ validation set
train_df, val_df = train_test_split(train_df, test_size=0.2, random_state=42)

# Set up ImageDataGenerator for normalization
train_datagen_Cfv = ImageDataGenerator(rescale=1./255) # Normalize pixel values_
↳ to [0, 1]
val_datagen_Cfv = ImageDataGenerator(rescale=1./255)
test_datagen_Cfv = ImageDataGenerator(rescale=1./255)

# Set up generators for training, validation, and test sets
print('For Training Image Set:')
```



```

train_generator_CFV = train_datagen_CFV.flow_from_dataframe(dataframe=train_df,
    ↪x_col='Filepath',
    y_col='Label',
    ↪target_size=(224, 224),
    batch_size=32,
    ↪class_mode='categorical')
print('\nFor Validation Image Set:')
val_generator_CFV = val_datagen_CFV.flow_from_dataframe(dataframe=val_df,
    ↪x_col='Filepath',
    y_col='Label',
    ↪target_size=(224, 224),
    batch_size=32,
    ↪class_mode='categorical')
print('\nFor Test Image Set:')
test_generator_CFV = test_datagen_CFV.flow_from_dataframe(dataframe=test_df,
    ↪x_col='Filepath', y_col='Label',
    target_size=(224, 224),
    ↪batch_size=32,
    class_mode='categorical',
    ↪shuffle=False)

```

For Training Image Set:

Found 1536 validated image filenames belonging to 4 classes.

For Validation Image Set:

Found 384 validated image filenames belonging to 4 classes.

For Test Image Set:

Found 480 validated image filenames belonging to 4 classes.

```

[111]: # -----
# -----
# IMAGE CLASSIFICATION WITH GENERAL MODEL CFV
# -----
# -----

# Define number of classes
Num_classes = 4
# Create the general CFV model
general_model_CFV = build_model(Num_classes)
# Compile Model
general_model_CFV.compile(
    optimizer=tf.keras.optimizers.SGD(learning_rate = 1e-3, momentum=0.9),
    loss='categorical_crossentropy',
    metrics=['accuracy'])

```

```

# Create KFold cross-validation object
kf = KFold(n_splits=5, shuffle=True, random_state=42)

# Perform cross-validation
for train_index, val_index in kf.split(train_df):
    train_set, val_set = train_df.iloc[train_index], train_df.iloc[val_index]

    # Train model using train_generator and validate on val_generator
    general_model_CV_history = general_model_CV.fit(train_generator_CV,
    ↪ epochs=5, validation_data=val_generator_CV)

```

Epoch 1/5

```

2024-03-06 21:04:20.221746: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
[[{{node Placeholder/_0}}]]

```

48/48 [=====] - ETA: 0s - loss: 1.3028 - accuracy: 0.3971

```

2024-03-06 21:04:51.438658: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
[[{{node Placeholder/_0}}]]

```

48/48 [=====] - 33s 672ms/step - loss: 1.3028 - accuracy: 0.3971 - val_loss: 1.1855 - val_accuracy: 0.4531

Epoch 2/5

48/48 [=====] - 31s 650ms/step - loss: 1.0415 - accuracy: 0.5553 - val_loss: 1.1058 - val_accuracy: 0.5260

Epoch 3/5

48/48 [=====] - 31s 650ms/step - loss: 0.8891 - accuracy: 0.6465 - val_loss: 0.8289 - val_accuracy: 0.6536

Epoch 4/5

48/48 [=====] - 34s 696ms/step - loss: 0.7201 - accuracy: 0.7324 - val_loss: 0.8210 - val_accuracy: 0.6927

Epoch 5/5

48/48 [=====] - 32s 676ms/step - loss: 0.5687 - accuracy: 0.7923 - val_loss: 0.6554 - val_accuracy: 0.7448

Epoch 1/5

```

2024-03-06 21:07:09.057888: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
[[{{node Placeholder/_0}}]]

```

```

48/48 [=====] - ETA: 0s - loss: 0.4446 - accuracy:
0.8438

2024-03-06 21:07:40.927708: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
      [[{{node Placeholder/_0}}]]

48/48 [=====] - 33s 693ms/step - loss: 0.4446 -
accuracy: 0.8438 - val_loss: 0.5232 - val_accuracy: 0.7839
Epoch 2/5
48/48 [=====] - 34s 707ms/step - loss: 0.4051 -
accuracy: 0.8620 - val_loss: 0.4395 - val_accuracy: 0.8281
Epoch 3/5
48/48 [=====] - 32s 677ms/step - loss: 0.3050 -
accuracy: 0.8971 - val_loss: 0.3909 - val_accuracy: 0.8516
Epoch 4/5
48/48 [=====] - 32s 679ms/step - loss: 0.2999 -
accuracy: 0.8945 - val_loss: 0.3591 - val_accuracy: 0.8698
Epoch 5/5
48/48 [=====] - 32s 662ms/step - loss: 0.2436 -
accuracy: 0.9193 - val_loss: 0.3367 - val_accuracy: 0.8672
Epoch 1/5

2024-03-06 21:10:00.719636: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
      [[{{node Placeholder/_0}}]]

48/48 [=====] - ETA: 0s - loss: 0.1898 - accuracy:
0.9368

2024-03-06 21:10:30.124989: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
      [[{{node Placeholder/_0}}]]

48/48 [=====] - 31s 645ms/step - loss: 0.1898 -
accuracy: 0.9368 - val_loss: 0.3427 - val_accuracy: 0.8828
Epoch 2/5
48/48 [=====] - 33s 680ms/step - loss: 0.1692 -
accuracy: 0.9466 - val_loss: 0.3265 - val_accuracy: 0.8724
Epoch 3/5
48/48 [=====] - 31s 639ms/step - loss: 0.1488 -
accuracy: 0.9512 - val_loss: 0.3620 - val_accuracy: 0.8542
Epoch 4/5
48/48 [=====] - 31s 638ms/step - loss: 0.1316 -
accuracy: 0.9590 - val_loss: 0.2912 - val_accuracy: 0.8932

```

```

Epoch 5/5
48/48 [=====] - 33s 684ms/step - loss: 0.0942 -
accuracy: 0.9720 - val_loss: 0.2905 - val_accuracy: 0.8828
Epoch 1/5

2024-03-06 21:13:07.128373: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
      [[{{node Placeholder/_0}}]]

48/48 [=====] - ETA: 0s - loss: 0.0744 - accuracy:
0.9831

2024-03-06 21:13:37.207336: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
      [[{{node Placeholder/_0}}]]

48/48 [=====] - 31s 656ms/step - loss: 0.0744 -
accuracy: 0.9831 - val_loss: 0.2648 - val_accuracy: 0.8880
Epoch 2/5
48/48 [=====] - 32s 662ms/step - loss: 0.0600 -
accuracy: 0.9870 - val_loss: 0.2700 - val_accuracy: 0.8984
Epoch 3/5
48/48 [=====] - 31s 638ms/step - loss: 0.0516 -
accuracy: 0.9915 - val_loss: 0.2913 - val_accuracy: 0.8958
Epoch 4/5
48/48 [=====] - 33s 682ms/step - loss: 0.0377 -
accuracy: 0.9941 - val_loss: 0.3217 - val_accuracy: 0.8906
Epoch 5/5
48/48 [=====] - 31s 644ms/step - loss: 0.0563 -
accuracy: 0.9837 - val_loss: 0.3508 - val_accuracy: 0.8854
Epoch 1/5

2024-03-06 21:15:44.417507: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
      [[{{node Placeholder/_0}}]]

48/48 [=====] - ETA: 0s - loss: 0.0563 - accuracy:
0.9831

2024-03-06 21:16:14.998507: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
      [[{{node Placeholder/_0}}]]

```

```

48/48 [=====] - 32s 672ms/step - loss: 0.0563 -
accuracy: 0.9831 - val_loss: 0.3100 - val_accuracy: 0.9036
Epoch 2/5
48/48 [=====] - 31s 647ms/step - loss: 0.0295 -
accuracy: 0.9987 - val_loss: 0.3033 - val_accuracy: 0.8984
Epoch 3/5
48/48 [=====] - 30s 620ms/step - loss: 0.0224 -
accuracy: 0.9980 - val_loss: 0.2887 - val_accuracy: 0.9089
Epoch 4/5
48/48 [=====] - 31s 654ms/step - loss: 0.0163 -
accuracy: 0.9987 - val_loss: 0.3648 - val_accuracy: 0.8854
Epoch 5/5
48/48 [=====] - 33s 695ms/step - loss: 0.0297 -
accuracy: 0.9948 - val_loss: 0.2999 - val_accuracy: 0.9089

```


```

[112]: plt.figure(figsize=(10, 4))
plt.subplot(1, 2, 1)
plt.plot(general_model_CFV_history.history['loss'], label='Training Loss')
plt.plot(general_model_CFV_history.history['val_loss'], label='Validation_
↳Loss')
plt.title('General Model CFV Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(general_model_CFV_history.history['accuracy'], label='Training_
↳Accuracy')
plt.plot(general_model_CFV_history.history['val_accuracy'],
↳label='Validation Accuracy')
plt.title('General Model CFV Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.tight_layout();
plt.show();

```



```
[113]: # Evaluate the model on the test set
general_model_CFV_predictions = general_model_CFV.predict(test_generator)
```


2024-03-06 21:46:13.459184: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID_ARGUMENT: You must feed a value for placeholder tensor 'Placeholder/_0' with dtype int32

[[{{node Placeholder/_0}}]]

15/15 [=====] - 2s 124ms/step

```
[129]: # Create confusion matrix
general_model_CFV_predicted_labels = np.argmax(general_model_CFV_predictions,
↪axis=1)
true_labels = test_generator.classes
conf_matrix = confusion_matrix(true_labels, general_model_CFV_predicted_labels)

# Plot the confusion matrix
plt.figure(figsize=(5, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Purples',
↪xticklabels=test_generator.class_indices, yticklabels=test_generator.
↪class_indices)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix -- General Model CFV')
plt.show();
```



```
[115]: # Count mislabeled images
# Get class names from folder names
class_names = sorted(os.listdir(dataset_path))

# Get true labels
true_labels = test_generator.classes

# Get predicted labels
general_model_CFV_predicted_labels = np.argmax(general_model_CFV_predictions,
↪axis=1)

# Find mislabeled indices
general_model_CFV_mislabeled_indices = np.where(true_labels !=
↪general_model_CFV_predicted_labels)[0]

print("Number of mislabeled images:", len(general_model_CFV_mislabeled_indices))
print("Number of predictions:", len(general_model_CFV_predictions))
print("Overall accuracy:", round((1- (len(general_model_CFV_mislabeled_indices)/
↪len(general_model_CFV_predictions)))*100, 2), "percent" )
```

Number of mislabeled images: 13
 Number of predictions: 480
 Overall accuracy: 97.29 percent

```
[116]: # -----
# -----
# IMAGE CLASSIFICATION WITH GOOGLNET MODEL CFV
# -----
# -----

# Create the googlenet CFV model
googlenet_model_CFV = googlenet()
# Compile Model
googlenet_model_CFV.compile(
    optimizer=tf.keras.optimizers.SGD(learning_rate = 1e-3, momentum=0.9),
    loss='categorical_crossentropy',
    metrics=['accuracy'])

# Create KFold cross-validation object
kf = KFold(n_splits=5, shuffle=True, random_state=42)

# Perform cross-validation
for train_index, val_index in kf.split(train_df):
    train_set, val_set = train_df.iloc[train_index], train_df.iloc[val_index]

    # Train model using train_generator and validate on val_generator
    googlenet_model_CFV_history = googlenet_model_CFV.fit(train_generator_CFV,
↳ epochs=5, validation_data=val_generator_CFV)
```

Epoch 1/5

```
2024-03-06 21:47:32.957328: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
[[{{node Placeholder/_0}}]]
```

```
48/48 [=====] - ETA: 0s - loss: 1.3613 - accuracy:
0.3587
```

```
2024-03-06 21:53:44.064873: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
[[{{node Placeholder/_0}}]]
```

```
48/48 [=====] - 399s 8s/step - loss: 1.3613 - accuracy:
0.3587 - val_loss: 1.3175 - val_accuracy: 0.5703
```

Epoch 2/5

```
48/48 [=====] - 397s 8s/step - loss: 1.2228 - accuracy:
0.5534 - val_loss: 1.1703 - val_accuracy: 0.4688
```

Epoch 3/5

```
48/48 [=====] - 398s 8s/step - loss: 1.0347 - accuracy:
```



```

0.5853 - val_loss: 0.9466 - val_accuracy: 0.6016
Epoch 4/5
48/48 [=====] - 398s 8s/step - loss: 0.8494 - accuracy:
0.6686 - val_loss: 0.8986 - val_accuracy: 0.6068
Epoch 5/5
48/48 [=====] - 394s 8s/step - loss: 0.7842 - accuracy:
0.7018 - val_loss: 0.8068 - val_accuracy: 0.6562
Epoch 1/5

2024-03-06 22:22:52.176890: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
      [[{{node Placeholder/_0}}]]

48/48 [=====] - ETA: 0s - loss: 0.7019 - accuracy:
0.7409

2024-03-06 22:29:00.630462: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
      [[{{node Placeholder/_0}}]]

48/48 [=====] - 397s 8s/step - loss: 0.7019 - accuracy:
0.7409 - val_loss: 0.7218 - val_accuracy: 0.7135
Epoch 2/5
48/48 [=====] - 391s 8s/step - loss: 0.5930 - accuracy:
0.7754 - val_loss: 0.6926 - val_accuracy: 0.7188
Epoch 3/5
48/48 [=====] - 391s 8s/step - loss: 0.5368 - accuracy:
0.8092 - val_loss: 0.5970 - val_accuracy: 0.7734
Epoch 4/5
48/48 [=====] - 394s 8s/step - loss: 0.4764 - accuracy:
0.8333 - val_loss: 0.5413 - val_accuracy: 0.8177
Epoch 5/5
48/48 [=====] - 396s 8s/step - loss: 0.4557 - accuracy:
0.8464 - val_loss: 0.4995 - val_accuracy: 0.8359
Epoch 1/5

2024-03-06 22:58:05.283045: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
      [[{{node Placeholder/_0}}]]

48/48 [=====] - ETA: 0s - loss: 0.4217 - accuracy:
0.8444

2024-03-06 23:04:07.184683: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an

```

error and you can ignore this message): INVALID_ARGUMENT: You must feed a value for placeholder tensor 'Placeholder/_0' with dtype int32

[[{{node Placeholder/_0}}]]

48/48 [=====] - 390s 8s/step - loss: 0.4217 - accuracy: 0.8444 - val_loss: 0.4810 - val_accuracy: 0.8255

Epoch 2/5

48/48 [=====] - 395s 8s/step - loss: 0.3898 - accuracy: 0.8646 - val_loss: 0.5265 - val_accuracy: 0.7969

Epoch 3/5

48/48 [=====] - 395s 8s/step - loss: 0.3509 - accuracy: 0.8841 - val_loss: 0.4515 - val_accuracy: 0.8385

Epoch 4/5

48/48 [=====] - 394s 8s/step - loss: 0.3556 - accuracy: 0.8763 - val_loss: 0.4831 - val_accuracy: 0.8099

Epoch 5/5

48/48 [=====] - 396s 8s/step - loss: 0.3281 - accuracy: 0.8848 - val_loss: 0.4138 - val_accuracy: 0.8411

Epoch 1/5

2024-03-06 23:34:09.206017: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID_ARGUMENT: You must feed a value for placeholder tensor 'Placeholder/_0' with dtype int32

[[{{node Placeholder/_0}}]]

48/48 [=====] - ETA: 0s - loss: 0.3631 - accuracy: 0.8672

2024-03-06 23:40:16.021625: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID_ARGUMENT: You must feed a value for placeholder tensor 'Placeholder/_0' with dtype int32

[[{{node Placeholder/_0}}]]

48/48 [=====] - 396s 8s/step - loss: 0.3631 - accuracy: 0.8672 - val_loss: 0.3982 - val_accuracy: 0.8333

Epoch 2/5

48/48 [=====] - 393s 8s/step - loss: 0.3186 - accuracy: 0.8828 - val_loss: 0.4197 - val_accuracy: 0.8464

Epoch 3/5

48/48 [=====] - 395s 8s/step - loss: 0.2723 - accuracy: 0.9062 - val_loss: 0.3921 - val_accuracy: 0.8516

Epoch 4/5

48/48 [=====] - 408s 8s/step - loss: 0.2504 - accuracy: 0.9128 - val_loss: 0.3848 - val_accuracy: 0.8490

Epoch 5/5

48/48 [=====] - 406s 8s/step - loss: 0.2481 - accuracy: 0.9167 - val_loss: 0.4311 - val_accuracy: 0.8359

Epoch 1/5

```
2024-03-07 00:09:49.679834: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
```

```
[[{{node Placeholder/_0}}]]
```

```
48/48 [=====] - ETA: 0s - loss: 0.2325 - accuracy:
0.9186
```

```
2024-03-07 00:16:02.251971: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'Placeholder/_0' with dtype int32
```

```
[[{{node Placeholder/_0}}]]
```

```
48/48 [=====] - 401s 8s/step - loss: 0.2325 - accuracy:
0.9186 - val_loss: 0.3606 - val_accuracy: 0.8568
```

```
Epoch 2/5
```

```
48/48 [=====] - 396s 8s/step - loss: 0.2185 - accuracy:
0.9264 - val_loss: 0.3893 - val_accuracy: 0.8516
```

```
Epoch 3/5
```

```
48/48 [=====] - 393s 8s/step - loss: 0.2054 - accuracy:
0.9264 - val_loss: 0.3670 - val_accuracy: 0.8698
```

```
Epoch 4/5
```

```
48/48 [=====] - 391s 8s/step - loss: 0.2167 - accuracy:
0.9238 - val_loss: 0.3323 - val_accuracy: 0.8802
```


```
Epoch 5/5
```

```
48/48 [=====] - 391s 8s/step - loss: 0.1988 - accuracy:
0.9316 - val_loss: 0.3496 - val_accuracy: 0.8750
```

```
[117]: plt.figure(figsize=(10, 4))
plt.subplot(1, 2, 1)
plt.plot(googlenet_model_CFV_history.history['loss'], label='Training Loss')
plt.plot(googlenet_model_CFV_history.history['val_loss'], label='Validation_
↳Loss')
plt.title('GoogleNet Model CFV Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(googlenet_model_CFV_history.history['accuracy'], label='Training_
↳Accuracy')
plt.plot(googlenet_model_CFV_history.history['val_accuracy'],
↳label='Validation Accuracy')
plt.title('GoogleNet Model CFV Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
```

```
plt.tight_layout();
plt.show();
```




```
[118]: # Evaluate the model on the test set
googlenet_model_CFV_predictions = googlenet_model_CFV.predict(test_generator)
```

2024-03-07 00:50:42.224899: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an error and you can ignore this message): INVALID_ARGUMENT: You must feed a value for placeholder tensor 'Placeholder/_0' with dtype int32
 [[{{node Placeholder/_0}}]]

15/15 [=====] - 35s 2s/step

```
[119]: # Create confusion matrix
googlenet_model_CFV_predicted_labels = np.
    ↳argmax(googlenet_model_CFV_predictions, axis=1)
true_labels = test_generator.classes
conf_matrix = confusion_matrix(true_labels,
    ↳googlenet_model_CFV_predicted_labels)

# Plot the confusion matrix
plt.figure(figsize=(5, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Greens',
    ↳xticklabels=test_generator.class_indices, yticklabels=test_generator.
    ↳class_indices)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('General Model Confusion Matrix')
plt.show();
```



```
[120]: # Count mislabeled images
# Get class names from folder names
class_names = sorted(os.listdir(dataset_path))

# Get true labels
true_labels = test_generator.classes

# Get predicted labels
googlenet_model_CFV_predicted_labels = np.
    ↳argmax(googlenet_model_CFV_predictions, axis=1)

# Find mislabeled indices
googlenet_model_CFV_mislabeled_indices = np.where(true_labels !=
    ↳googlenet_model_CFV_predicted_labels)[0]

print("Number of mislabeled images:",
    ↳len(googlenet_model_CFV_mislabeled_indices))
print("Number of predictions:", len(googlenet_model_CFV_predictions))
print("Overall accuracy:", round((1-
    ↳len(googlenet_model_CFV_mislabeled_indices)/
    ↳len(googlenet_model_CFV_predictions))*100, 2), "percent" )
```

Number of mislabeled images: 41

Number of predictions: 480
Overall accuracy: 91.46 percent

```
[131]: #  
#  
# SUMMARY OF PERFORMANCES FOR NON CROSSFOLD VALIDATION (CFV) GENERAL AND  
# GOOGLNET MODEL VARIATIONS  
# SUMMARY OF PERFORMANCES FOR CROSSFOLD VALIDATION (CFV) GENERAL AND GOOGLNET  
# MODEL VARIATIONS  
#  
#  
# General Model Confusion Matrix  
plt.figure(figsize=(15,10))  
plt.subplot(2, 2, 1)  
general_model_predicted_labels = np.argmax(general_model_predictions, axis=1)  
true_labels = test_generator.classes  
conf_matrix = confusion_matrix(true_labels, general_model_predicted_labels)  
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',  
            xticklabels=test_generator.class_indices, yticklabels=test_generator.  
            class_indices)  
plt.xlabel('Predicted Labels')  
plt.ylabel('True Labels')  
plt.title('Confusion Matrix -- General Model')  
  
# GoogleNet Model Confusion Matrix  
plt.subplot(2, 2, 2)  
google_model_predicted_labels = np.argmax(google_model_predictions, axis=1)  
true_labels = test_generator.classes  
conf_matrix = confusion_matrix(true_labels, google_model_predicted_labels)  
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Reds',  
            xticklabels=test_generator.class_indices, yticklabels=test_generator.  
            class_indices)  
plt.xlabel('Predicted Labels')  
plt.ylabel('True Labels')  
plt.title('Confusion Matrix -- GoogleNet Model')  
  
# General Model CFV Confusion Matrix  
plt.subplot(2, 2, 3)
```

```

general_model_CFV_predicted_labels = np.argmax(general_model_CFV_predictions,
↪axis=1)
true_labels = test_generator.classes
conf_matrix = confusion_matrix(true_labels, general_model_CFV_predicted_labels)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Purples',
↪xticklabels=test_generator.class_indices, yticklabels=test_generator.
↪class_indices)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix -- General Model CFV')

# GoogleNet Model CFV Confusion Matrix
plt.subplot(2, 2, 4)
googlenet_model_CFV_predicted_labels = np.
↪argmax(googlenet_model_CFV_predictions, axis=1)
true_labels = test_generator.classes
conf_matrix = confusion_matrix(true_labels,
↪googlenet_model_CFV_predicted_labels)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Greens',
↪xticklabels=test_generator.class_indices, yticklabels=test_generator.
↪class_indices)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix -- GoogleNet Model CFV')

plt.tight_layout();
plt.show();

general_accuracy = str(round((1- (len(general_mislabeled_indices)/
↪len(general_model_predictions)))*100, 2)) + ' %'
google_accuracy = str(round((1- (len(google_model_mislabeled_indices)/
↪len(google_model_predictions)))*100, 2)) + ' %'
general_CFV_accuracy = str(round((1- (len(general_model_CFV_mislabeled_indices)/
↪len(general_model_CFV_predictions)))*100, 2)) + ' %'
google_CFV_accuracy = str(round((1-
↪len(googlenet_model_CFV_mislabeled_indices)/
↪len(googlenet_model_CFV_predictions)))*100, 2)) + ' %'

# Organize Data
data = {'Model': ['General',
                  'GoogleNet',
                  'General CFV',
                  'GoogleNet CFV',
                  ],
        'Predictions': [len(general_model_predictions),
                         len(google_model_predictions),

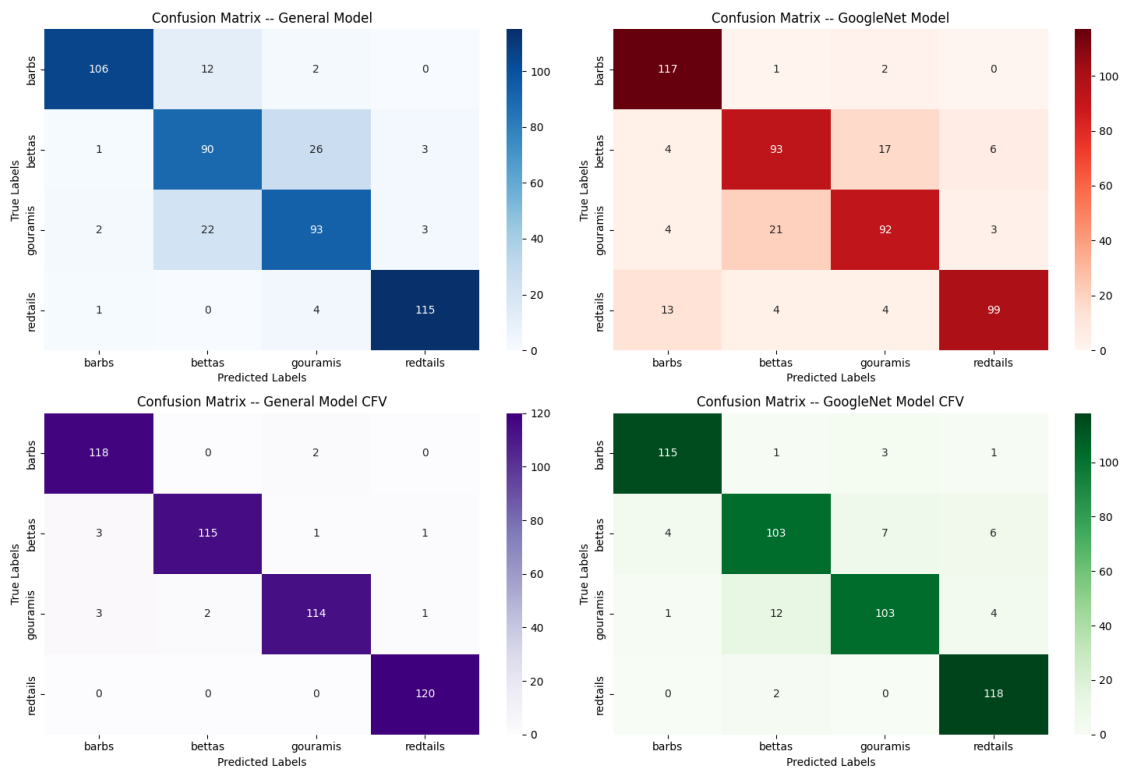
```

```

        len(general_model_CFV_predictions),
        len(googlenet_model_CFV_predictions),
    ],
    'Errors': [len(general_mislabeled_indices),
               len(google_model_mislabeled_indices),
               len(general_model_CFV_mislabeled_indices),
               len(googlenet_model_CFV_mislabeled_indices)
    ],
    'Accuracy': [general_accuracy,
                 google_accuracy,
                 general_CFV_accuracy,
                 google_CFV_accuracy,
    ]}

# Construct a DataFrame
df = pd.DataFrame.from_dict(data).set_index('Model')
# Display the DataFrame
print(df);

```



	Predictions	Errors	Accuracy
Model			
General	480	76	84.17 %
GoogleNet	480	79	83.54 %
General CFV	480	13	97.29 %

GoogleNet CFV 480 41 91.46 %

OBSERVATIONS FOR ALL MODEL VARIATIONS:


- 1– Using crossfold validation (CFV) improved both general and googlenet model accuracies
- 2– Even after CFV, the googlenet model still had trouble differentiating gouramis from bettas (12 errors)
- 3– The general model with CFV (“General Model CFV”) seems to be the best model
- 4– For both models, each row contains 120 image predictions as expected

```
[139]: # Display all mislabeled images with class names for model with best_
↳ performance ("General Model CFV")
plt.figure(figsize=(10,10))
for i in range(0, len(general_model_CFV_mislabeled_indices)):
    index = general_model_CFV_mislabeled_indices[i]
    true_label_index = true_labels[index]
    predicted_label_index = general_model_CFV_predicted_labels[index]


    img_path = test_generator.filepaths[index]
    img = tf.keras.preprocessing.image.load_img(img_path)
    img_array = tf.keras.preprocessing.image.img_to_array(img) / 255.0

    plt.subplot(4, 4, i + 1)
    plt.imshow(img)
    plt.title(f'True: {class_names[true_label_index]} \n Predicted:↳
↳ {class_names[predicted_label_index]}')
    plt.axis('off')
plt.tight_layout()
plt.show();
```


True: barbs
Predicted: gouramis



True: barbs
Predicted: gouramis



True: bettas
Predicted: redtails




True: bettas
Predicted: gouramis




True: bettas
Predicted: barbs



True: bettas
Predicted: barbs




True: bettas
Predicted: barbs




True: gouramis
Predicted: barbs




True: gouramis
Predicted: barbs



True: gouramis
Predicted: barbs



True: gouramis
Predicted: bettas



True: gouramis
Predicted: bettas



True: gouramis
Predicted: redtails

